

Retrieval-Augmented Generation: From Simple to Advanced Systems

Introduction to RAG

Retrieval-Augmented Generation, commonly known as RAG, represents one of the most significant advancements in the field of natural language processing and large language models. At its core, RAG is a framework that enhances the capabilities of generative models by incorporating external knowledge retrieval mechanisms. This approach addresses one of the fundamental limitations of traditional language models: their reliance on static training data and the subsequent inability to access or incorporate information beyond their training cutoff.

The basic premise of RAG is elegantly simple yet profoundly powerful. Instead of relying solely on the parametric knowledge stored within a model's weights, RAG systems first retrieve relevant information from an external knowledge source and then use this retrieved information to augment the generation process. This combination of retrieval and generation allows for more accurate, up-to-date, and contextually relevant responses.

Chapter 1: The Foundations of RAG

The Problem with Traditional Language Models

Traditional language models, including early transformer-based architectures, operate purely on the knowledge they acquired during training. Once trained, these models have no mechanism to access new

information or verify facts against external sources. This creates several significant limitations:

First, the knowledge becomes stale over time. A model trained in 2021 has no awareness of events, discoveries, or developments that occurred after its training cutoff. Second, these models have no inherent mechanism for fact verification—they generate responses based on statistical patterns in their training data, which can lead to confident but incorrect statements, a phenomenon often called "hallucination." Third, the knowledge is implicit and distributed across billions of parameters, making it difficult to trace the source of information or update specific facts without retraining the entire model.

The Retrieval-Augmented Solution

RAG addresses these limitations by decoupling knowledge storage from generation capabilities. The generative model no longer needs to memorize all possible information; instead, it can focus on its core competency of understanding context and generating coherent text. The knowledge is stored externally in a retrievable format, typically a vector database or search index, which can be updated independently of the generative model.

This architectural separation offers several advantages. Knowledge can be updated simply by modifying the external database, without any retraining of the generative model. The retrieval step provides provenance—the system can cite sources and allow users to verify information. Additionally, the retrieval mechanism can be designed to access domain-specific

knowledge bases, enabling the same generative model to serve specialized applications in medicine, law, or other fields.

Chapter 2: Simple RAG Systems

The Basic Architecture

A simple RAG system consists of three primary components: a retrieval mechanism, a knowledge base, and a generative model. The retrieval mechanism converts queries into a format suitable for searching the knowledge base. The knowledge base contains documents or passages that have been pre-processed and indexed. The generative model takes the original query along with the retrieved documents and produces a final response.

The retrieval process typically begins with embedding both the query and the documents in the knowledge base into a shared vector space. When a user submits a query, it is converted into an embedding vector, and the system searches for documents whose embeddings are closest to this query embedding. This similarity search returns the most relevant documents, which are then passed to the generative model along with the original query.

Implementation Considerations

In simple RAG implementations, several design choices significantly impact performance. The choice of embedding model determines how effectively semantic relationships are captured. Early systems often used pre-trained models like

Sentence-BERT or Instructor embeddings. The chunking strategy—how documents are divided into retrievable units—affects both retrieval accuracy and the amount of context provided to the generator. Common approaches include fixed-size chunks with overlap, sentence-level chunks, or paragraph-level chunks.

The generative model in simple RAG systems is typically a large language model that has been fine-tuned to effectively utilize retrieved context. The prompt template must instruct the model to base its response primarily on the provided context while acknowledging when the context lacks sufficient information to answer the query.

Limitations of Simple RAG

Despite its effectiveness, simple RAG has notable limitations. The retrieval is typically performed once at query time, with no opportunity to refine or iterate based on the initial results. The retrieved documents are treated as independent pieces of information without considering their relationships or potential contradictions. The system has no mechanism for determining whether it has retrieved enough information or whether it needs to perform additional searches.

Chapter 3: Intermediate RAG Concepts

Chunking Strategies and Document Processing

As practitioners gained experience with RAG systems, it became clear that naive chunking strategies often suboptimal. The

ideal chunk size depends on multiple factors: the nature of the content, the typical query patterns, and the context window limitations of the generative model.

Small chunks provide precise retrieval but may lack sufficient context for the generator to understand the information properly. Large chunks provide more context but may contain irrelevant information that dilutes the relevance signal during retrieval. Advanced chunking strategies emerged, including recursive chunking that maintains document structure, semantic chunking that respects natural boundaries, and sliding window approaches that provide overlapping context.

Metadata and Filtering

Intermediate RAG systems incorporate metadata to enhance retrieval quality. Documents can be tagged with source information, publication dates, author credentials, document types, and domain classifications. During retrieval, this metadata enables filtering—for example, restricting results to documents published within a specific timeframe or from authoritative sources only.

Metadata also enables faceted search capabilities, where users can specify constraints on the retrieval process. In enterprise applications, this might include filtering by department, project, or document confidentiality level. The integration of metadata filtering with semantic search represents a significant advancement over simple vector similarity alone.

Hybrid Search Approaches

Vector similarity search excels at capturing semantic relationships but can miss exact matches or keyword-based relevance. Hybrid search combines vector search with traditional keyword matching (BM25, TF-IDF) to achieve the benefits of both approaches.

In hybrid search systems, the retrieval score combines a semantic similarity component from vector embeddings with a lexical matching component from keyword search. The weighting between these components can be tuned based on the application domain. For technical documents where precise terminology matters, keyword matching might receive higher weight. For more conceptual queries, semantic similarity might dominate.

Re-ranking for Improved Relevance

One of the most effective enhancements to simple RAG is the introduction of a re-ranking stage. Initial retrieval typically uses efficient but approximate methods to identify a candidate set of documents. These candidates are then passed through a more sophisticated but computationally expensive re-ranker that provides more accurate relevance scores.

Re-rankers are typically cross-encoders that process both the query and document together, allowing for deeper interaction between the texts. While too slow for initial retrieval over large corpora, they excel at refining a smaller candidate set. This two-stage approach combines the scalability of bi-encoder retrieval with the accuracy of cross-encoder re-ranking.

Chapter 4: Advanced RAG Architectures

Multi-Query Retrieval

Advanced RAG systems recognize that a single query may not adequately capture the user's information need. Multi-query retrieval generates multiple variations of the original query, each emphasizing different aspects or using different phrasings. These query variations are used to retrieve documents, and the results are aggregated to provide a more comprehensive set of relevant information.

The query variations can be generated by the language model itself, which can identify different facets of a complex question. For example, a query about "the impact of climate change on agriculture" might generate sub-queries about crop yield impacts, regional variations, adaptation strategies, and economic consequences. This approach ensures that the retrieval covers multiple dimensions of the topic.

Self-Reflective RAG

Self-reflective RAG systems incorporate mechanisms for the system to evaluate its own retrieval and generation quality. After retrieving documents and generating a response, the system reflects on whether the retrieved documents actually support the generated answer, whether any critical information might be missing, and whether the response fully addresses the query.

This reflection can trigger additional retrieval steps. If the system detects that key information is missing or that the generated response contains unsupported

claims, it can formulate new queries to fill these gaps. This iterative process continues until the system determines that it has sufficient information or reaches a maximum iteration limit.

Hypothetical Document Embeddings (HyDE)

HyDE represents an innovative approach to bridging the vocabulary gap between queries and documents. Instead of embedding the query directly, HyDE first uses a language model to generate a hypothetical document that would answer the query. This hypothetical document is then embedded and used for retrieval.

The intuition behind HyDE is that documents similar to the hypothetical document are likely to contain relevant information. This approach is particularly effective when queries are short or when there's a significant vocabulary mismatch between how questions are asked and how information is presented in documents. The hypothetical document serves as a prototype of what a relevant document might look like.

Contextual Compression

As retrieved document sets grow, they often contain redundant information or passages that are only peripherally relevant. Contextual compression addresses this by extracting or summarizing the most relevant portions of retrieved documents before passing them to the generator.

Compression can take several forms. Extractive compression selects the most

relevant sentences or passages from each document. Abstractive compression generates concise summaries of longer documents, focusing on aspects relevant to the query. The compressed context allows the generator to focus on the most pertinent information while fitting within context window limitations.

Chapter 5: Evaluation and Optimization

Retrieval Metrics

Evaluating RAG systems requires assessing both retrieval quality and generation quality. For retrieval, standard information retrieval metrics apply, including precision at k ($P@k$), recall at k , mean average precision (MAP), and normalized discounted cumulative gain (NDCG). These metrics measure how well the retrieval system surfaces relevant documents.

However, retrieval relevance in RAG systems has a unique characteristic: documents are relevant to the extent that they help the generator produce a good response. This suggests that retrieval evaluation should consider downstream utility, not just topical relevance. A document might be topically relevant but contain outdated information that would lead to an incorrect response.

Generation Quality Assessment

Evaluating generated responses presents additional challenges. Traditional metrics like BLEU and ROUGE measure overlap with reference responses but correlate poorly with human judgments of quality. More sophisticated approaches include:

Factual consistency metrics that verify whether generated claims are supported by retrieved documents. Answer relevance metrics that assess whether the response addresses the query. Completeness metrics that evaluate whether the response covers all aspects of the query. Fluency and coherence metrics that assess the quality of the generated text.

Human evaluation remains the gold standard but is expensive and difficult to scale. Recent approaches use strong language models as judges, prompting them to evaluate responses along multiple dimensions. While not perfect, LLM-based evaluation provides a scalable approximation of human judgment.

Latency and Cost Optimization

Production RAG systems must balance quality with practical constraints of latency and cost. Retrieval from large vector databases can become a bottleneck, especially when using exact nearest neighbor search. Approximate nearest neighbor (ANN) algorithms provide dramatic speed improvements with minimal accuracy degradation.

Caching strategies can reduce redundant computations. Frequently asked queries and their responses can be cached. Document embeddings can be pre-computed and stored. The generative model's context processing can be optimized through techniques like prefix caching.

Cost considerations often drive architectural decisions. The generative model typically represents the largest

computational cost. Strategies to reduce this cost include using smaller models for straightforward queries, reserving larger models for complex cases, and optimizing prompt lengths to minimize token usage.

Chapter 6: Advanced Techniques and Future Directions

Graph-Based RAG

Recent innovations incorporate graph structures into RAG systems. Instead of treating documents as independent chunks, graph-based RAG captures relationships between entities, documents, and concepts. Knowledge graphs provide structured representations that complement the unstructured text in traditional RAG.

In graph RAG systems, retrieval operates over both text chunks and graph nodes. A query might retrieve relevant entities from the knowledge graph, then traverse graph relationships to find related information. This approach excels at multi-hop reasoning tasks that require connecting information from multiple sources.

Graph structures also enable more sophisticated context assembly. Retrieved information can be organized according to graph relationships, providing the generator with not just relevant text but also an understanding of how pieces of information relate to each other.

Multimodal RAG

As language models evolve to handle multiple modalities, RAG systems are expanding beyond text. Multimodal RAG can retrieve and incorporate images, diagrams, tables, audio, and video alongside text. This is particularly valuable in domains where information is inherently multimodal, such as medical imaging with radiology reports or technical documentation with diagrams.

Multimodal retrieval requires embedding all modalities into a shared representation space. Recent advances in multimodal models like CLIP provide exactly this capability—text and images can be embedded such that similar concepts have similar vector representations regardless of modality.

The generation component must also handle multiple modalities, either by generating multimodal responses or by effectively incorporating retrieved multimodal context into text generation. This remains an active area of research with rapid progress.

Agentic RAG

The frontier of RAG development involves integrating agent capabilities. Rather than passive retrieval followed by generation, agentic RAG systems can take multiple actions, observe results, and adapt their strategies. These systems maintain state across interactions, remembering what information they've already retrieved and what gaps remain.

Agentic RAG can use tools beyond simple retrieval—they might execute code to analyze data, query APIs for real-time information, or interact with databases

through structured queries. The agent decides what actions to take based on the current state and the ultimate goal of answering the user's query.

This approach enables handling of complex tasks that require multiple steps and diverse information sources. For example, answering "What was the revenue trend for our competitors over the last five years and how does it correlate with their R&D spending?" might require retrieving financial reports, extracting relevant numbers, performing calculations, and synthesizing insights.

Personalization and User Modeling

Advanced RAG systems increasingly incorporate user context and personalization. The same query might require different responses depending on the user's expertise level, role, or prior knowledge. Personalization can be implemented at multiple stages of the RAG pipeline.


During retrieval, user context can influence relevance scoring. A query about "neural networks" should retrieve different documents for a high school student versus a machine learning researcher. During generation, user context influences the explanation depth, technical vocabulary, and examples provided.

User modeling extends beyond static profiles to include interaction history. The system remembers what information a user has already seen and can avoid redundancy while building on previously established understanding. This creates a more natural and efficient interaction pattern.

Chapter 7: Practical Implementation Considerations

Choosing the Right Components

Building a RAG system requires numerous architectural decisions. The embedding model selection depends on the language domain, whether the content is domain-specific or general, and the desired balance between retrieval quality and computational efficiency. Popular options range from small, efficient models like all-MiniLM-L6-v2 to larger, more capable models like OpenAI's text-embedding-3-large.



The vector database choice involves tradeoffs between scalability, feature set, and operational complexity. Options include purpose-built vector databases like Pinecone or Weaviate, vector search plugins for existing databases like pgvector for PostgreSQL, and cloud-native services like Amazon Kendra.

The generative model selection must consider capability requirements, latency constraints, cost considerations, and whether deployment will be cloud-based or

on-premises. Organizations must also decide between using commercial APIs or self-hosting open-source models.

Managing Knowledge Updates

A key advantage of RAG is the ability to update knowledge without model retraining. However, implementing effective knowledge updates requires careful consideration. When new documents are added, they must be processed, chunked, embedded, and indexed. When documents are updated or removed, the index must reflect these changes.

For dynamic knowledge bases, update frequency becomes important. Some applications require real-time updates as new information arrives. Others can tolerate batch updates on a daily or weekly schedule. The update mechanism must maintain consistency—users should not see partial or inconsistent information during update processes.

Versioning and rollback capabilities become important in regulated industries. The ability to reconstruct what knowledge was available at a specific point in time may be required for audit or compliance purposes.

Handling Edge Cases

Robust RAG systems must handle various edge cases gracefully. When no relevant documents are retrieved, the system should acknowledge its limitations rather than attempting to answer from its parametric knowledge alone. When retrieved documents contain contradictory information, the system should identify the

contradiction and possibly request clarification.

Multi-lingual scenarios require embedding models and generative models that support all relevant languages. Cross-lingual retrieval—queries in one language retrieving documents in another—may be desired in some applications but requires careful consideration of embedding model capabilities.

Domain adaptation often requires fine-tuning components on domain-specific data. While the basic RAG architecture works across domains, optimal performance typically requires domain-specific tuning of chunking strategies, embedding models, and prompt templates.

Chapter 8: Case Studies and Applications

Enterprise Knowledge Management

Large organizations have adopted RAG systems to make their internal knowledge more accessible. Technical documentation, project reports, meeting notes, and institutional knowledge stored in wikis and document repositories become queryable through natural language interfaces.

In these applications, access control becomes critical. The retrieval system must respect document permissions—users should only retrieve documents they're authorized to see. This adds complexity to the retrieval pipeline, requiring integration with identity management systems and filtering based on user permissions.

Customer Support Automation

RAG-powered customer support systems combine product documentation, support ticket histories, and knowledge base articles to provide accurate, contextually relevant responses to customer inquiries. These systems reduce support costs while improving consistency and availability.

Advanced implementations integrate with customer relationship management (CRM) systems to personalize responses based on customer history, product usage, and subscription tier. The retrieval system can prioritize information relevant to the customer's specific situation.

Research and Analysis

In research-intensive domains, RAG systems help analysts and researchers navigate vast document collections. Legal research, scientific literature review, and competitive intelligence analysis all benefit from RAG's ability to surface relevant information from large corpora.

These applications often require sophisticated query understanding to handle complex information needs. Researchers might need to find documents that support or contradict specific hypotheses, track the evolution of ideas over time, or identify relationships between concepts across document collections.

Education and Learning

Educational applications of RAG provide personalized learning experiences.

Students can ask questions about course materials and receive explanations tailored to their current understanding. The system can retrieve relevant sections from textbooks, supplementary materials, and example problems.

Adaptive learning systems use RAG to assemble personalized study materials. Based on a student's demonstrated knowledge gaps, the system retrieves and presents appropriate explanations and practice problems, creating a customized learning path.

Chapter 9: Challenges and Limitations

Factual Accuracy and Hallucination

Despite significant advances, RAG systems still struggle with factual accuracy. Retrieved documents may contain incorrect information, and the generative model may misinterpret or misrepresent what it retrieves. The system may also combine information from multiple sources in ways that create false implications not present in any individual source.

Mitigation strategies include rigorous source verification, confidence scoring for generated claims, and explicit citation of sources. However, complete elimination of factual errors remains an open challenge, particularly for complex queries requiring synthesis across multiple documents.

Context Window Limitations

Even with expanding context windows in modern language models, the amount of

information that can be provided to the generator remains limited. When relevant information exceeds these limits, the system must make difficult decisions about what to include and what to omit. This compression inevitably loses potentially valuable information.

Strategies to address context limitations include hierarchical summarization, where documents are summarized and then summaries are summarized, and dynamic context selection that adapts the provided context based on the generation progress. These approaches help but don't fully solve the fundamental constraint.

Bias and Fairness

RAG systems can amplify and perpetuate biases present in their knowledge bases. If the retrieved documents reflect historical biases or underrepresent certain perspectives, the generated responses will reflect these same biases. The retrieval mechanism itself may introduce additional biases through its relevance judgments.

Addressing bias requires careful curation of knowledge bases, diverse representation in training data for all components, and ongoing monitoring of system outputs for fairness concerns. Techniques like fairness-aware retrieval and debiasing during generation are active research areas.

Computational and Economic Costs

Production RAG systems incur significant computational costs. Embedding large document collections, maintaining vector indexes, and running generative models all

require substantial infrastructure. For applications at scale, these costs can become prohibitive.

Cost optimization strategies include tiered retrieval (cheaper methods first, more expensive only when needed), model quantization to reduce memory and computation requirements, and strategic caching of common queries and their results. Organizations must carefully balance quality requirements against operational costs.

Chapter 10: The Future of RAG

Integration with Foundation Models

As foundation models continue to evolve, the boundary between retrieval and generation is blurring. Some models now incorporate retrieval capabilities directly into their architecture, learning when and how to retrieve information as part of their training. These "retrieval-augmented language models" may eventually make the distinction between parametric and non-parametric knowledge less relevant.

Long-context models capable of processing entire books may reduce the need for sophisticated chunking and retrieval in some applications. However, the computational efficiency of retrieval over exhaustive context processing will likely ensure that RAG remains valuable for very large knowledge bases.

Reasoning and Multi-Step Capabilities

Future RAG systems will incorporate more sophisticated reasoning capabilities. Rather than simple retrieval-then-generate, they will engage in multi-step reasoning, retrieving information, forming hypotheses, retrieving additional information to test these hypotheses, and iteratively refining their understanding.

This evolution toward reasoning systems will require closer integration with symbolic AI approaches, combining neural retrieval and generation with explicit reasoning over structured knowledge. The result will be systems that don't just retrieve and parrot information but genuinely reason with it.

Personalization and Adaptation

RAG systems will become increasingly personalized, learning from interactions with individual users to better understand their knowledge, preferences, and communication style. This personalization will extend beyond surface-level adaptation to fundamentally shape what information is retrieved and how it's presented.

Privacy-preserving personalization techniques will be essential, allowing systems to learn from user interactions without compromising sensitive information. Federated learning and on-device personalization may play important roles.

Ethical and Governance Frameworks

As RAG systems become more capable and widely deployed, robust governance frameworks will be essential. These frameworks must address questions of

accountability when systems provide incorrect information, transparency about information sources, and mechanisms for correcting errors or removing inappropriate content.

Standards for evaluating RAG systems across different dimensions—accuracy, fairness, robustness, transparency—will emerge, enabling meaningful comparisons and certifications. Regulatory frameworks may establish requirements for certain high-stakes applications.

Conclusion

Retrieval-Augmented Generation represents a fundamental shift in how we build systems that interact with knowledge. By combining the flexibility and fluency of large language models with the precision and updatability of retrieval systems, RAG addresses many of the limitations that constrained earlier approaches.

The journey from simple to advanced RAG systems has been marked by continuous innovation—better retrieval techniques, more sophisticated integration with generation, and growing understanding of how to evaluate and optimize these systems. Yet the field remains in its early stages, with many challenges still to address and opportunities yet to explore.

For practitioners and researchers, RAG offers a rich landscape for innovation. The basic architecture admits countless variations and enhancements, each suited to different applications and constraints. As the field evolves, we can expect RAG systems to become more capable, more

reliable, and more seamlessly integrated into how we work with information.

The future of RAG is not just about better retrieval or better generation in isolation, but about creating systems that genuinely understand and work with knowledge in ways that augment human intelligence. This vision—of AI systems that can access, reason with, and communicate knowledge effectively—is what continues to drive progress in this exciting field.

Field	Available from Docling?
File name, file path	Yes (but you already have this — you passed it in)
Number of pages	Yes
Document size	Yes (but you get this from the filesystem, not Docling)

Challenge:

Mehrfach gedacht, besser entschieden: Prio-1-Projekte mit Self-Consistency Prompting

Organisationen bearbeiten parallel eine Vielzahl interner Projekte, die um Ressourcen, Aufmerksamkeit und Priorität konkurrieren. Oft ist unklar, welches Projekt die höchste Priorität haben sollte, da verschiedene Kriterien (z. B. Relevanz oder Nutzen) abgewogen werden müssen. GenKI kann diesen Entscheidungsprozess unterstützen, indem sie Abwägungen transparent macht und verschiedene Argumentationslinien durchspielt. Eine dafür geeignete Methode ist Self-Consistency Prompting, bei dem die GenKI mehrfach unabhängig argumentiert und dann die Ergebnisse miteinander verglichen werden, um konsistente Muster zu erkennen.

HINTERGRUND

AUFGABE

Erstellt mithilfe der GenKI eine priorisierte Liste interner Projekte. Ausgangspunkt ist eine bestehende oder selbst erstellte Liste laufender/geplanter interner Projekte. Zunächst formuliert einen Prompt, der die GenKI dazu anleitet, die Priorisierung mehrfach unabhängig vorzunehmen (in verschiedenen unabhängigen Chats). Die GenKI soll dabei jedes Mal begründen, welches Projekt sie als höchste Priorität einstuft und warum. Anschließend vergleicht die unterschiedlichen Ergebnisse, identifiziert wiederkehrende Muster und prüft, welches Projekt über die verschiedenen Durchläufe hinweg am häufigsten als Prio-1-Projekt erscheint. Abschließend reflektiert:

Inwiefern führt Self-Consistency Prompting zu besseren Priorisierungsentscheidungen?

Nutzt selbst erstellte Beispiele oder sucht euch aus dem Internet/Intranet passende heraus.

MATERIAL